

CHRISTOPHER J. HIGGINS, Ph.D.

AI Safety & Evaluation | Product Policy

Brooklyn, NY | ch.ja.higgins@gmail.com | [LinkedIn](#) | [Substack](#) | [Washington Blade](#) | [Academic Papers](#)

PROFILE

LLM evaluator and analytical writer focused on reasoning quality, product policy, and model behavior. I build and review LLM benchmarks, rubrics, QA guidance, and adversarial prompts; assess outputs for logical coherence, ambiguity, hidden assumptions, evidentiary fit, and calibration; and publish public analysis on AI behavior, anthropomorphism, chatbot intimacy, AI welfare, and LLM citation norms.

CORE FIT

AI policy & safety | Policy analysis & operationalization | Risk assessment | Adversarial simulation & red teaming | Responsible AI | Model behavior analysis | Policy guidance & documentation | QA & calibration | Cross-functional communication | Prompt engineering | Python (light)

EXPERIENCE

AI Safety & Red Teaming Consultant | July AI | Remote | 2026-present

- Design safety-focused red-team scenarios that probe policy boundaries across hate speech, minors, violence and graphic content, refusals and over-refusals, and other sensitive-content edge cases.
- Evaluate model outputs for policy adherence, classification consistency, factual rigor, calibration, and user-safety failures; document recurring failure modes and actionable feedback.

LLM Benchmark Team Lead | Turing | Remote | 2025-2026

- Led a domain team building LLM benchmarks; translated ambiguous capability goals into operational criteria, reviewer guidance, and QA briefs.
- Reviewed contributor work for clarity, consistency, logical coherence, measurability, and rubric fit; diagnosed recurring model-behavior and evaluation failure modes.

LLM Benchmark Contributor | Scale AI / Outlier AI - Humanity's Last Exam | Remote | 2024-2025

- Authored doctoral-level prompts, structured rubrics, and formal reasoning tasks for a high-stakes benchmark assessing frontier LLM capabilities.
- Evaluated model answers for ambiguity, logical validity, unstated assumptions, reasoning errors, and failure to satisfy rubric criteria in a double-review workflow.

Author & Editor | Curdled Incompleteness Theorem | Remote | 2025-present

- Publish essays on model behavior, AI anthropomorphism, chatbot intimacy, AI welfare, LLM citation practices, and conversational-AI design.
- Translate technical and philosophical disputes into public analysis by separating architectural claims, behavioral evidence, and conceptual interpretation; portfolio reaches 2k+ monthly views across 26 countries.

Feature Writer | Washington Blade | Washington, DC | 2023-2024

- Reported features on politics, law, and culture; conducted interviews and turned primary-source material into analytical narratives for a national audience; 67% selected as cover stories.

Postdoctoral Fellow and Instructor | Johns Hopkins University | Baltimore, MD | 2016-2024

- Published peer-reviewed research on Wittgenstein and ordinary language methodology; taught writing-intensive courses in political theory and conversation theory.
- Trained students to reconstruct arguments, identify hidden assumptions, steel-man rival interpretations, and weigh competing evidentiary standards.

SELECTED RESEARCH & WRITING

- [California's Crackdown on Chatbot Therapy for Minors](#)- model-behavior experiment on chatbot intimacy, safety prompts, boundary-setting, and minor protection.
- [Why Graph Theory Lost the AI Race](#)- technical essay on graph theory, linear algebra, transformer architecture, and why GPTs model language through vectors rather than explicit relations.
- [SEC Data Collection Project](#)- workflow-design case study translating sparse specs into worker tiers, task categories, automation, QA rubrics, revision loops, and delivery estimates.
- [Cautions about \(Cautions about\) AI Ethics](#)- analysis of AI welfare, moral uncertainty, and how ethical caution can become institutional authority.

EDUCATION & TRAINING

Ph.D., Political Science (Political Theory), Johns Hopkins University, 2023

B.A., Mathematics and Political Science, Williams College, 2014

IBM Generative AI Engineering coursework: Python, prompt engineering, data analysis, machine learning.