

CHRISTOPHER J. HIGGINS, Ph.D.

AI Evaluation | Research | Technical Communication

Brooklyn, NY | ch.ja.higgins@gmail.com | [linkedin.com/in/higginscj](https://www.linkedin.com/in/higginscj) | [cj-higgins.github.io](https://github.com/cj-higgins)

PROFILE

AI evaluator, researcher, and analytical writer with a Ph.D. in political theory and a B.A. in mathematics. Design model-behavior evaluations, translate ambiguous goals into operational criteria, analyze quantitative and qualitative evidence, and communicate findings clearly to technical and nontechnical audiences. Experience spans frontier-model red-teaming, benchmark QA, original empirical research, workflow design, journalism, and public technical writing.

CORE CAPABILITIES

Model-behavior evaluation & red-teaming | Evaluation design | Rubric and QA systems | Quantitative analysis | Metric development | Research synthesis | Workflow design | Technical and public-facing writing | Stakeholder-ready reporting | Python, pandas, Jupyter (working proficiency)

EXPERIENCE

AI Safety Evaluation & Red-Teaming Consultant | July AI | Remote | Mar 2026-present

- Design and evaluate adversarial conversations probing safety-policy boundaries, refusals and over-refusals, sensitive-content handling, and ambiguous edge-case model behavior for frontier AI labs.
- Document recurring failure patterns and translate evolving guidance into consistent, deliverable-ready findings during compressed sprint cycles.

LLM Benchmark Team Lead | Turing | Remote | 2025-2026

- Led political-science benchmark work on a four-person domain team, reviewing hundreds of prompts and evaluations within a multi-pass QA process combining expert review, automated grader checks, divergence analysis, and targeted follow-up.
- Developed a taxonomy of benchmark defects, calculated error rates, and wrote a stakeholder memo identifying ambiguous answer keys, flawed questions, and question-answer mismatches; the client used the findings to commission a replacement benchmark.

Author & Editor | Curdled Incompleteness Theorem | Remote | 2025-present

- Own an independent publication on AI, technology, and model behavior; pitch, research, write, edit, and publish original analysis reaching 2k+ monthly views across 26 countries.
- Designed original studies of roughly 14,000 chatbot-customer-service interactions and 324 LLM responses; built comparison metrics, used logistic regression, and created charts to surface product and policy implications.
- Translate fast-moving AI developments into clear arguments by separating behavioral evidence, technical architecture, product incentives, and conceptual claims.

LLM Benchmark Contributor | Scale AI / Outlier AI - Humanity's Last Exam | Remote | 2024-2025

- Authored doctoral-level prompts, structured rubrics, and formal reasoning tasks for a high-stakes benchmark assessing frontier-model capabilities.
- Evaluated model answers for ambiguity, logical validity, unstated assumptions, reasoning errors, and failure to satisfy rubric criteria in a double-review workflow.

Feature Writer | Washington Blade | Washington, DC | 2023-2024

- Reported bimonthly features on politics, law, and culture; conducted interviews and primary-source research and shaped source-heavy material into publication-ready narratives under deadline.
- Identified the central implication in complex institutional and policy stories for a national audience; 67% of reported features were selected as cover stories.

Postdoctoral Fellow & Instructor | Johns Hopkins University | Baltimore, MD | 2016-2024

- Published peer-reviewed research, delivered research talks, and designed and taught four writing-intensive courses translating complex theory into structured explanations for varied audiences.
- Earned a 4.75 teaching-effectiveness score against a 4.14 university average; trained students to analyze evidence, compare interpretations, and communicate concise conclusions.

SELECTED WORK

- **Trusting the Chatbot More Than the Front Desk** - Quantitative analysis of roughly 14,000 customer-service dialogues; developed a sentiment measure and validated it with logistic regression to compare direct escalation with behind-the-scenes human consultation.
- **California's Crackdown on Chatbot Therapy for Minors** - A 324-response GPT-4o-mini study comparing baseline and policy-prompted behavior across warmth, boundary-setting, jailbreak resistance, and possible policy violations.
- **SEC Filing Data Collection Workflow** - End-to-end workflow for producing 100 question-answer pairs from SEC 10-K filings, with automation, human review, QA standards, revision loops, and delivery tracking.
- **Why Graph Theory Lost the AI Race** - Technical explainer connecting graph neural networks, transformer architecture, language representation, sparse matrices, and GPU hardware.

EDUCATION & TRAINING

Ph.D., Political Science (Political Theory), Johns Hopkins University, 2023

B.A., Mathematics and Political Science, Williams College, 2014

IBM Generative AI Engineering coursework: Python, pandas, prompt engineering, data analysis, and machine learning.