


From Generic Travel Landing Pages to Experience-Specific Design

A Contra Labs case study

Diagnostic findings and a contributor-trainable plan for a 1,000-pair dataset



Two phases—and the delivery timeline

TWO PHASES

1

Phase 1 · Diagnostic

Evaluate the model's outputs and identify the most important failure.

2

Phase 2 · Dataset


Build 1,000 reviewed prompt → ideal pairs designed to fix it.

Timeline

Sequence gates, staffing, economics, and delivery.

High-consideration guided travel

Boutique multi-day tour booking · narrow enough to measure, valuable enough to matter



● Commercial stakes

High-consideration purchases depend on trust, clarity, and evidence of delivery—not just atmosphere.

● Stable conventions

Trip pages share understandable jobs while still allowing legitimate variation in layout, voice, and imagery.

● Plausible staffing

Contra category pages show a credible sourcing path for web and brand designers; pilot throughput still needs proof.

Diagnostic evaluation

Turn vague “AI-generated” complaints into measurable page-level and set-level failure.

Different trips, same opening grammar



01 Tuscany



02 Japan Family



03 Iceland Photography



04 Morocco*



05 Costa Rica




06 Namibia


* Morocco: two intended CSS background images were blocked by Chromium ORB. Treat the empty fields as a technical caveat, not art-direction evidence.

Tuscany and Iceland repeat the same hero grammar

TUSCANY · FIRST PASS



ICELAND · FIRST PASS



Different destinations, identical art-direction grammar: abstract mark · atmospheric fade · serif statement + colored italics.

The pages work; but the experiences are not fully tangible

Six first-pass outputs · frozen 1440 × 900 renders · rubric locked, then applied retrospectively
Directional mini-eval · AI-assisted scoring informed by qualitative human review; production evaluation uses calibrated human raters.

PAGE-LEVEL MEAN · 0–2 SCALE



SET-LEVEL VARIATION

- 1 / 2 Layout
- 0 / 2 Typography
- 1 / 2 Photography
- 0 / 2 Voice

HOW TO READ THIS

Each criterion contains observable sub-criteria scored 0–2. Bars show the mean across six pages; set-level variation is scored separately.

Phase 1 Operations

01 · SOURCE

Find credible design judges

Contra web, product, and brand designers.

02 · CALIBRATE

Measure the bench

Each candidate scores one fresh page and one heterogeneous set.

03 · SCORE + CONTROL

Run paired expert judgment

Every page and set batch gets two blind ratings. Material 0↔2 disagreement triggers review.

BASE-CASE PLANNING · LOCKED AFTER CALIBRATION

DEFAULT STAFFING

4 evaluators

USEFUL MIDDLE RANGE

120–150 outputs

CREATIVE SPEND

\$6.0K–\$7.35K

INDICATIVE TIMING


≈1.5–2 weeks

Fine-tuning dataset

Build 1,000 reviewed prompt → ideal pairs around one specific, learnable mechanism.

Failure Mode: the trip is clearer than the experience

JAPAN · FIRST PASS · ACTIVITY MODULE



SELECTED FAILURE

Visual evidence adapts too little to the brief

The destination changes more readily than what the page visually proves. Even when participation, instruction, craft, or hosted delivery carries the value, the imagery often remains generic.

PAGE SIGNAL

Experience Tangibility · 1.67 / 2

People & Participation: 1 / 2 on 5 of 6 pages


SET SIGNAL

Photography Variation · 1 / 2

TRAINING OBJECTIVE

Infer the product value → choose imagery that proves it.

Phase 2 Operations



BASE-CASE PLANNING · LOCKED AFTER THE PAID PILOT



Prevent one house style from taking over the dataset

BUILD VARIATION INTO THE DATASET

- **Plan the mix** — Spread briefs across trip types, travelers, brands, price points, and evidence needs.
- **Cross the assignments** — No designer repeatedly owns the same kind of brief.
- **Break the scaffold** — Designers own the final page and can reshape or discard the AI starting point.

CATCH COLLAPSE DURING PRODUCTION

- **Audit batches as a set** — Compare pages across designers and brief types for repeated layout, typography, imagery, and voice.
- **Find the source** — If one pattern takes over, check the brief mix, creator assignments, and AI starting point.
- **Correct before scaling** — Rebalance the work, rework affected pages, and verify the fix before continuing production.

Japan: a reviewed prototype of a higher design bar

BEFORE · FIRST PASS



AFTER · REVIEWED PROTOTYPE




Taiko, walking, cycling, and a logistics ledger make participation and hosted delivery visible.

Example 1 of 3 · two additional prototypes in Appendix G · validation remains a pilot gate.

Project timeline

Sequence gates, staffing, throughput, risk, and delivery across ten weeks.

A ten-week plan



Main operating risk: the pilot reveals slower throughput, weaker agreement, or more rework.
Mitigation: add calibrated capacity or extend schedule—don’t compromise quality.

Blocking dependencies: model interface compatibility · rights-cleared assets · client approval cadence

Two budgets, staged commercial locks

PHASE 1 · DIAGNOSTIC

120–150 outputs

recommended planning middle

\$6.0K

creative spend at 120

\$7.35K

creative spend at 150

\$6.0K–\$7.35K is planned creative payout within the ≈\$10K ceiling. Final N locks after calibration and may flex if other billable delivery costs must fit.

PHASE 2 · DATASET

\$175K–\$225K

fixed-fee range · planning case ≈ \$200K

1,000

accepted pairs

≈200

steady-state
pairs / week

14

10 designers · 3
reviewers · 1 lead

≈50 base review hrs/week + up to ≈10 reserve hrs. \$116.25K base labor / \$120K full reserve. Planning range now; final fixed fee locks after the paid pilot validates throughput, rejection, and rework.

What we need from you before the pilot

Five answers that shape how we build and deliver the dataset.

CLIENT QUESTION

WHAT IT CHANGES

1

What should each training example include?



Defines exactly what we produce, review, and deliver.

2

What technical constraints must every page follow?



Sets the framework, styling rules, accessibility requirements, and tooling we build around.

3

What can the model actually see about the images it uses?



Determines whether we can train the visual-evidence problem directly—or need to change the training target.

4


How will you use this dataset?



Determines what assets we can use, what metadata we need, and how we separate training from evaluation data.

5

Who will review the work, and how quickly can they respond?



Sets the feedback loop, rework process, staffing plan, and timeline.